



WARWICK BUSINESS SCHOOL
THE UNIVERSITY OF WARWICK

Student Number	5667249
Module Code	IB9YC0
Module Title	Data Driven Decision Making
Submission Deadline	22-Jan-2026 12:00:00 PM
Date Submitted	22-Jan-2026 11:49:29 AM
Word Count	2366
Number of Pages	21
Question Attempted	All the given questions from the question paper are solved according to the given order.
Have you used Artificial Intelligence (AI) in any part of this assignment?	Yes
If you have ticked “Yes” above, please briefly outline below which AI tool you have used, and what you have used it for. Please note, you must also reference the use of generative AI correctly within your assessment, in line with the guidance provided in your student handbook. I used generative AI tools (ChatGPT and Google Gemini) solely to generate referee-style evaluations of the assigned research papers, as explicitly required by the assessment instructions. These AI-generated evaluations were used as comparative benchmarks against my own independent evaluations.	

Table of Contents

Section 1: Paper One - Are Economics Students Biased Against Female Teachers? Evidence from a Randomized, Double-Blind Natural Field Experiment.....	3
1.1. My Thoughts.....	3
1.2. ChatGPT vs Gemini.....	4
1.3. Human Judgement vs LLM Evaluations	4
Section 2: Paper 2 - Mission Motivation and Public Sector Performance: Experimental Evidence from Pakistan	6
2.1. My Thoughts.....	6
2.2. ChatGPT vs Gemini	7
2.3. Human Judgement vs LLM Evaluation	7
Section 3:.....	9
3.1. LLMs as Research Paper Evaluators: Capability or Convenience?	9
3.2. LLMs Should Advice, Not Decide	9
References.....	10
Appendices.....	11

	ChatGPT's Score	Gemini's Score	My Score
Paper 1 Are Economics Students Biased Against Female Teachers? Evidence from a Randomized, Double-Blind Natural Field Experiment	78	88	80
Paper 2 Mission Motivation and Public Sector Performance: Experimental Evidence from Pakistan	85	94	82

Section 1: Paper One - Are Economics Students Biased Against Female Teachers? Evidence from a Randomized, Double-Blind Natural Field Experiment

1.1. My Thoughts

This paper explores the hypothesis of how student evaluations of teachings (SETs) are discriminatory against female instructors. To answer this question, the authors provide a randomized and double-blind natural field experiment in a university setting. Students were encouraged to send course-related questions to a generic email address, and the reply would be signed either a male or a female name. Importantly, even the instructors were not aware of the gendered identity assigned to each response. Student rated instructors based on helpfulness, knowledge, and speed of response after the course. This design permits the authors to casually approximate the impact of perceived instructor gender on student evaluations.

The paper's contribution is primarily methodological. By implementing a double-blind technique, it improves on earlier experimental studies that may have been affected by experimenter demand effects. The topic is clearly significant as SETs are now widely used in academic decision-making, and the paper is clearly written and well-structured, although it is technically demanding.

The experiment is well designed in terms of internal validity. The gender assignment is random, hypotheses are preregistered, and balance tests show no systematic differences between treatment conditions. Several estimation strategies are applied, and the authors effectively prove that null results are not due to low statistical power. These aspects help to argue that the paper has validly defined the casual influence of perceived gender in this context.

However, the study's external validity is more limited. The context is very specific: online, written communication in a Scandinavian university with a rather high student satisfaction. These conditions may reduce the expression of bias and restrict generalizability to in-person teachings or other institutional settings. From a behavioral science perspective, the paper also relies on a reduced-form approach, concluding the absence of bias from outcomes without directly testing psychological mechanisms such as stereotype activation or expectation effects.

To refine the work, future experiments could be conducted to experimentally manipulate the strength of gender cues or manipulate contextual factors, which may induce bias. Interpretation could also be improved by a more intensive adherence to the behavioral theory. The paper philosophically indicates that there was a positive focus on causal identification,

rather than theory development, which makes it more credible but less insightful in understanding behavior.

1.2. ChatGPT vs Gemini

The two analyses of the LLM have quite similar strong aspects but differ in the weaknesses that they focus on and the grades they assign. Both reviews are quite categorical that the article covers a significant, as well as, policy-driven research question, and that possible gender bias in Student Evaluations of Teaching could have far-reaching implications on academic careers and gender equality. The same can be pointed out as the central strength of the paper and the primary source of novelty by both. More specifically, they concur that the blinding of the instructors to their perceived gender is a definite methodological improvement on previous studies and a significant improvement in internal validity. Additionally, both the LLMs take the null finding to be a good one, stating that the narrow confidence intervals in the paper are capable of ruling out the effect sizes found in highly influential previous research.

The reviews however differ in their determination of weaknesses. The presentation in the paper is also more critiqued in the ChatGPT's evaluation which states that the partial-derivatives explanation and the verbatim emails correspondence are ill-placed within the manuscript and it undermines the professionalism of the manuscript unless more properly framed or summarised. It also conveys the expression of caution about the analysis aspect that cannot be seen in the excerpt completely. On the contrary, these aspects are not perceived as weaknesses in Gemini's evaluation. Rather, it applauds the inclusion of original correspondence as increasing transparency and having strong confidence in the empirical analysis and estimation approach, and considers the limitations of external validity as having been taken into account.

The scores are therefore not comparable: in ChatGPT's evaluation, 78/100 is scored, in Gemini's evaluation, 88/100 is scored, which is different but not an agreement concerning the essence of the contribution made by the paper.

1.3. Human Judgement vs LLM Evaluations

The broad features of my assessment are consistent with the LLM assessments in determining the key strengths of the paper, specifically, the significance of the research question and the methodological value of the double-blind randomized field experiment. Similar to the LLMs, I also highlight that instructor blinding is more effective in enhancing internal validity and is also objectionably superior to previous experimental designs. I also subscribe to the positive interpretation of the null result, in agreement with the fact that the authors effectively prove that low statistical power is not a factor that causes the absence of bias.

Nevertheless, my consideration favours more external validity and behaviour interpretation over either of the two analyses in LLM. Although the LLMs do recognize contextual constraints, they do not regard them as the major priority to a large extent. On the other hand, I perceive the very specific environment, which is online and text-based communication in one of the Scandinavian universities with high levels of student satisfaction, as a substantive limitation to generalisability to in-person teaching settings. I further point out that the paper is reduced-form and does not test psychological processes directly like stereotype activation which restricts its behavioural understanding. The focus of the LLM evaluations where identification and execution is a priority eliminates these issues mostly.

On the other hand, on the weaknesses associated with presentation, including the incorporation of technical explanations and verbatim email messages, ChatGPT evaluation ranks higher than I do. Although I say that the paper is technically demanding, the stylistic issues do not mean so much to me as causal credibility.

These discrepancies are the cause of these differences in scores. I would probably have an intermediate result between the two scores of LLM indicating a high degree of confidence in internal validity but a more stringent opinion towards generalisability and theoretical depth.

Section 2: Paper 2 - Mission Motivation and Public Sector Performance: Experimental Evidence from Pakistan

2.1. My Thoughts

The paper examines how the focus on organizational mission can enhance the performance of the employees working in the public sector within the low-income communities. Using an experimental intervention in Pakistan's public health system, the authors examine whether mission-oriented messaging can motivate workers after they have already joined the organization, rather than merely influencing worker-job matching. The study focuses on community health workers and changes in behavior like household visits and service delivery.

The article answers a very topical question in behavioral science and development economics, which connects intrinsic motivation, organizational design, and delivery of the services in the public. This contribution is made in a setting that obviously has a gap in the literature on the continuation of motivation, but not recruitment. The article is well composed and well organized on a general basis and the paper has a narrative that is well linked regarding theory, intervention, and outcomes.

A major strength is the experimental design. The study can be used to draw credible causal inference that the mission emphasis has on worker behavior by using a mission-based treatment and comparing it to financial incentives and placebo conditions. Strong internal validity is supported by randomization and using administrative and survey data and the findings are convincing in demonstrating that the mission treatment resulted in meaningful behavioral change compared to the status quo.

The external validity of the paper, however, is reduced. The sample, as it is identified by the authors, only comprises of female workers in one district which is relatively high on the Human Development Index in Pakistan. These contextual attributes lead to the question of whether the same effect would be noticed in male workers, in smaller capacity areas or on other workplaces in the public sector. Additionally, while the mission treatment improved performance, overall output remained far below policy targets, suggesting that motivation alone cannot address deeper structural constraints.

From a behavioral science perspective, further attempts to unpack the psychological processes involved in the paper would have enhanced it. Although the argument focuses on intrinsic motivation, it is not directly measured but deduced. The next step in research would be to integrate mission messaging with mission identity, mission norms or pro-social preferences measures to further understand why behavior changes. Philosophically, the

research is reflexive of a positivist, interventionist paradigm concerned more with causal identification than theory building, which increases policy relevance but reduces theory.

2.2. ChatGPT vs Gemini

The two LLMs reviews have many similarities in their strengths, though the relative strengths vary and in general assessment. Both concur that the research question, whether stressing organization mission can enhance performances of incumbent public sector workers is very crucial and policy relevant in development and public economics. Both point to poor quality of delivering public services as a fundamental issue in low-income nations and identify that the paper provides field-experimental evidence of great value to the literature that had been dominated by theories and laboratory experiments previously. They also agree that the novelty of the paper is an empirical approach and not a conceptual one.

Nonetheless, the reviews differ in the degree in which they frame the contribution of the paper. ChatGPT's tone is more reserved, stating that the contribution is sound and relevant to policy but has a chance to be narrowed down to discuss more specifically and elaborate more on the mechanisms or heterogeneity. Conversely, the evaluation by Gemini seems to make a unique and possibly ground-breaking contribution, especially by moving the perspective of worker selection to motivation of existing workers and by thoroughly investigating the interplay between mission motivation and financial rewards. Gemini also emphasizes more on the sophistication of methods such as multiple arms of the treatment, use of placebos, multiple-testing corrections and strong checks.

The identified weaknesses are not completely overlapping and thus, ChatGPT and Gemini share rather superficial problems with presentation and analytical depth and presentational recommendations respectively. In keeping with this, the scores cannot be compared. The difference in judgement between the execution of the paper is 85/100 in case of ChatGPT and 94/100 in case of Gemini.

2.3. Human Judgement vs LLM Evaluation

My analysis is consistent with that of the LLM evaluations in terms of determining the major strengths of the paper. Like both LLMs, I attach relevance to the research question and policy significance and understand that the primary contribution of this paper is enabling credible field- experimental evidence relating to the topic of mission motivation among existing public sector employees. Also, I concur with the LLMs that the compelling point is the use of the experimental design, especially the comparison of the mission-based incentive, financial incentive, and placebo conditions, which allow making a strong causal inference, and internal validity.

My analysis, however, gives higher emphasis to external validity and behavioral interpretation as compared to both LLM analyses. Although the LLMs take into consideration contextual constraints, they consider them to mostly be secondary. Conversely, I believe that the limitation to female employees in a rather high-HDI district is a substantive limitation that poses significant questions of the generalizability of the results in terms of gender, geographic area, and public-sector position. I also underline that, despite some measurable improvements, the overall performance is still way below policy levels, which implies that motivation is inadequate to break the wider structural limitations. The LLM evaluations are mostly devoid of this point.

In the view of behavioral sciences, I also see a gap that the LLMs do not focus on sufficiently: intrinsic motivation is not measured but rather assumed. Though the exploration of mechanisms is highly appreciated in the context of Gemini, I believe that the direct measurements of mission identity, mission norms, or pro-social preferences could be used to help clarify the motivation behind behavior changes. On the other hand, the LLMs, especially Gemini, have more emphasis on the sophistication and practice of methodology than I do, which helps them score higher. My evaluation would be most probably between the two scores of LLM, which would present strong confidence with causal identification, but more cautiousness with respect to theoretical depth and generalizability.

Section 3:

3.1. LLMs as Research Paper Evaluators: Capability or Convenience?

Based on the three evaluations conducted for each paper, LLMs have a partial and structured ability to assess research papers but not in a comprehensive manner. In both scenarios, the LLMs have mentioned the key strengths that were consistent, namely the significance of the research question, the relevance of the policy, and the methodological rigor of the research. They were especially good when it comes to evaluation of internal validity, experimental design, and placing of contributions among the existing literature. In these aspects, their judgments were very close to academics and my judgment.

Nevertheless, the comparisons have shortcomings as well. Although there is consensus on the strengths, there were significant variations in the scores of the LLMs to which weaknesses were not weighted equally. More specifically, the external validity, theoretical depth, and mechanisms of behavior were underemphasized by LLMs, which is a dimension that needs contextual and disciplinary assessment. This means that, although structured methodological evaluation can be done using LLMs, broader interpretive and theoretical elements of research cannot be evaluated using the same.

3.2. LLMs Should Advice, Not Decide

Managers should use LLMs as decision support tools but not decision makers. LLMs are ideal in preliminary screening, contribution summarization and methodological consistency verification, particularly in cases where assessing a large body of research is involved. Nevertheless, ultimate decisions, especially those that are generalizable, theoretical and strategic applicability, must be left to human professionals.

Combination of multiple LLM assessments with the human judgment is the best strategy. This minimizes biasing due to personal model framing and makes sure that significant contextual and theoretical factors are not ignored and hence more informed decision making.

References

Andersson, O., Backman, M., Bengtsson, N. and Engström, P., 2025. *Are economics students biased against female teachers? Evidence from a randomized, double-blind natural field experiment*. *Journal of Economic Behavior & Organization*. Available at:

<https://www.sciencedirect.com/science/article/pii/S0167268125002720>

Khan, Muhammad Yasir (2025) 'Mission motivation and public sector performance: experimental evidence from Pakistan', *American Economic Review*. Available at:

<https://www.aeaweb.org/articles?id=10.1257/aer.20231247&ArticleSearch%5Bwithin%5D%5Barticletitle%5D=1&ArticleSearch%5Bwithin%5D%5Barticleabstract%5D=1&ArticleSearch%5Bq%5D=experiment&JelClass%5Bvalue%5D=0&journal=1&from=j>

OpenAI (2024) *ChatGPT* [Large language model]. Available at: <https://chat.openai.com/> (Accessed: 22 January 2026).

Google (2024) *Gemini* [Large language model]. Available at: <https://gemini.google.com/> (Accessed: 22 January 2026).

Appendices:

Paper 1

ChatGPT Output

i) Interest, importance, and originality of the research question

The research question—whether student evaluations of teaching are biased against female instructors—is clearly of high interest and importance. SETs are widely used in hiring, promotion, and teaching evaluation decisions, and evidence of systematic bias has substantial implications for gender equality in academia and the efficient allocation of talent. The topic is therefore well aligned with the interests of an economics journal, particularly in labour economics, education economics, and the economics of discrimination.

In terms of originality, the paper's main contribution lies in its experimental design. The use of a randomized, double-blind setup in a natural educational environment is a meaningful advancement over both observational studies and earlier audit-style experiments. In particular, the fact that instructors themselves are unaware of the gender identity under which they are perceived by students is novel and directly addresses concerns about differential effort, behaviour, or communication styles across treatment arms. This design improves internal validity relative to influential prior work such as MacNeill et al. (2015).

That said, the null finding—no detectable bias against female mentors—will likely be seen as controversial given the existing literature. While null results are valuable, especially when they credibly rule out effect sizes reported in earlier studies, the paper's contribution will ultimately depend on how convincingly it explains differences relative to prior findings (e.g., context, post-Covid environment, nature of online interaction, or student population). Overall, the research question is important and the design is sufficiently original to merit serious consideration.

Assessment: High interest and importance; good originality driven by design rather than topic novelty.

ii) Clarity and quality of writing and presentation

The abstract is clear, concise, and informative. It clearly states the motivation, experimental design, key novelty, and main findings, and it situates the results relative to influential prior studies. This is a strong point of the paper.

The introduction (based on the provided portion) is generally well written and well referenced. It clearly motivates the study, summarizes the relevant literature, and explains why experimental evidence is needed. The transition to prior experimental work is logical, though the introduction appears to cut off mid-sentence, making it difficult to assess how effectively the authors position their contribution relative to MacNell et al. (2015) and related studies.

However, the appended end material is problematic from a presentation standpoint. The inclusion of a detailed pedagogical explanation of partial derivatives, followed by what appears to be a verbatim email exchange between a student and a mentor, is highly unusual for an economics journal article. While this material may be relevant as an example of mentor–student interaction or as part of the experimental context, it is currently presented in a way that lacks framing, integration, and academic tone. As written, it reads more like course material or an appendix accidentally pasted into the manuscript rather than a polished section of a research paper.

The paper would benefit greatly from:

- Clear separation between main text and any illustrative or supplementary material.
- Rewriting or summarizing the email exchange in an anonymized, academic style, if it is indeed central to the analysis.
- Explicit explanation of why this material is included and how it relates to the empirical strategy or results.

Assessment: Core sections (abstract and introduction) are clear and professional; overall presentation quality is weakened by poorly integrated and stylistically inconsistent end material.

iii) Correctness and appropriateness of the analysis performed

Based on the description, the experimental design is appropriate and well suited to the research question. Random assignment of perceived instructor gender, combined with instructor blinding, directly targets the main identification challenge in this literature. The outcome measures (helpfulness, knowledge, response time) are standard and relevant dimensions of teaching evaluations.

The claim that confidence intervals exclude the effect sizes found in prior influential studies is particularly valuable, as it allows readers to assess whether the null result reflects low

power or a substantively different effect. Assuming the statistical implementation is sound (which cannot be fully verified from the excerpt), this is a strong analytical choice.

However, the provided material does not allow full evaluation of:

- Sample size and statistical power.
- Balance checks and randomization validity.
- Functional form and estimation methods.
- Robustness checks or heterogeneity analyses (e.g., by student gender, course type, or prior performance).

The appended explanation of partial derivatives does not contribute to the analysis and raises concerns about editorial oversight rather than econometric correctness. If this material is part of the experimental interaction, it should be analyzed systematically (e.g., tone, length, complexity) rather than presented descriptively.

Assessment: The analytical approach appears appropriate and potentially strong, but the paper's credibility depends on details not visible here and on better integration of illustrative material.

Overall evaluation and score

This paper addresses an important and policy-relevant question using a creative and credible experimental design. Its main strength lies in the double-blind setup and its ability to speak directly to influential prior findings. However, the manuscript in its current form suffers from presentation issues and incomplete or poorly integrated sections that detract from its professionalism and clarity.

Overall score: 78 / 100

This reflects a paper with strong potential and a solid core contribution, but one that requires substantial revision in presentation and clearer exposition of methods and results before being suitable for publication in a top economics journal.

Gemini Output

Title: Gender bias in student evaluations of teaching: Evidence from a double-blind randomized experiment Authors: O. Andersson et al.

I. Interest, Importance, and Originality of the Research Question

The research question—whether Student Evaluations of Teaching (SETs) are biased against female instructors—is of significant importance to the economics of education and labour economics. Given that SETs influence hiring, tenure, and promotion decisions, understanding the nature of potential discrimination (statistical vs. taste-based) is critical.

The originality of this paper lies in its methodological rigor. While previous literature has explored this topic, most studies are either observational (suffering from selection bias) or single-blind experiments. The authors correctly identify a critical flaw in the influential study by MacNeill et al. (2015): in previous audit studies, the instructors knew which gender they were impersonating, introducing the risk of "experimenter demand effects".

This study advances the literature significantly by implementing a double-blind field experiment. By creating an administrative intermediary to strip identifying metadata before the instructor responded, the authors ensured that the instructors were blind to their "assigned" gender. This effectively isolates student bias from potential changes in instructor behavior, a distinction often muddled in prior work.

II. Clarity and Quality of Writing and Presentation

The paper is written with high clarity and follows a logical structure. The experimental design is explained transparently, particularly the mechanism of the "email mentor" system and the randomization protocol.

- **Transparency:** The inclusion of actual email correspondence in the appendix (both Swedish originals and English translations) is excellent practice. It allows reviewers to verify that the interactions were "natural" and that the gendered signatures (Anton/Elin) were the primary treatment signal.
- **Data Visualization:** The timeline of emails (Fig 2) and the distribution of ratings (Fig 3) clearly illustrate the data flow and the lack of skewness in student responses.
- **Contextualization:** The authors do a good job of situating their study within the specific context of post-Covid hybrid teaching, acknowledging that this is a specific type of interaction (text-based, impersonal) compared to classroom teaching.

III. Correctness and Appropriateness of the Analysis Performed

The analysis is rigorous and appropriate for a Randomized Controlled Trial (RCT).

- **Experimental Design:** The use of pairwise block randomization is appropriate to ensure balance. The balance tests (Table 1) convincingly show that email characteristics (length, response time, attachments) were identical across the two perceived genders, validating the double-blind execution.
- **Estimation Strategy:** The authors employ a robust set of estimators: unadjusted means, OLS with controls (covs), and student fixed effects (fe). The use of student fixed effects is particularly strong as it controls for unobserved student heterogeneity, leveraging the fact that many students rated both mentors.
- **Power and Precision:** A key strength of the analysis is the discussion of statistical power. The authors note that their tight confidence intervals allow them to reject the large effect sizes reported in MacNeill et al. (2015). This transforms a "null result" into a meaningful finding: it suggests that previous large estimates of bias might be driven by design artifacts or specific contexts rather than a universal bias in all settings.
- **Robustness Checks:** The authors diligently checked for alternative mechanisms. They verified that the "null result" was not masking heterogeneity between male and female students and analysed "politeness" in follow-up emails using AI and human raters to ensure female mentors didn't face a more hostile environment.
- **Limitations:** The authors fairly acknowledge the "dosage" limitation—interaction via email is less intense than face-to-face teaching. They also discuss the external validity concern regarding the Scandinavian context, which has higher gender equality than the US context of previous studies.

IV. Overall Assessment

This is a high-quality execution of a field experiment that addresses a specific, identified weakness in the existing literature (lack of double-blind protocols). The "null result" is compelling because of the study's precise estimation and rigorous control of instructor behavior. While the setting (email mentoring) limits the ability to generalize fully to classroom lectures, the study provides a vital data point suggesting that gender bias in SETs may not be as ubiquitous or large in anonymous, text-based interactions as previously thought, provided the instructor's actual behavior is held constant.

Overall Score: 88/100

Paper 2

ChatGPT Output

Interest, importance, and originality of the research question

The paper addresses a highly important and policy-relevant question: whether emphasizing organizational mission can improve the performance of public sector workers in low-income settings. Weak public service delivery is a central challenge in development economics, and understanding low-cost, scalable ways to improve bureaucratic performance is of first-order importance. The focus on intrinsic motivation and mission orientation speaks directly to influential theoretical work (e.g., Besley and Ghatak) and complements a large empirical literature on incentives, monitoring, and governance.

The research question is not entirely new—mission motivation has been studied theoretically and in laboratory experiments—but the paper’s contribution lies in providing clean experimental evidence from a real-world public sector organization. The context (community health workers in Pakistan), the scale, and the field-experimental design substantially enhance the external relevance of the findings relative to prior lab-based studies. The paper also speaks to an active debate on whether “soft” management practices can meaningfully affect performance in bureaucracies, making it well suited for a leading economics or development journal.

Overall, the research question is important and timely, and the paper makes a clear original contribution by bringing credible field evidence to a literature that has largely lacked it.

Assessment: High interest and importance; solid originality through field experimentation rather than conceptual novelty.

ii) Clarity and quality of writing and presentation

The paper is generally well written, clearly structured, and accessible to a broad economics audience. The introduction effectively motivates the problem, situates the study within the existing literature, and clearly states the research question. The discussion of prior theoretical and experimental work is concise and relevant, and the contribution of the paper is easy to understand.

The exposition of the institutional context and experimental design is clear and detailed, which is particularly important given the policy setting. Results are presented in a

straightforward manner, and the concluding discussion thoughtfully reflects on external validity, limitations, and avenues for future research. The authors are commendably transparent about the constraints of the study (e.g., all-female sample, relatively high baseline human capital in the study area).

That said, the paper could benefit from some tightening in later sections, where discussion occasionally becomes repetitive or speculative. In particular, the concluding discussion of institutional inertia and managerial incentives, while interesting, could be more sharply linked back to the empirical findings. Nonetheless, these are relatively minor issues.

Assessment: Clear, professional, and well organized, with minor scope for concision and sharper focus in discussion sections.

iii) Correctness and appropriateness of the analysis performed

The empirical strategy is appropriate and well aligned with the research question. The use of a randomized field experiment provides a strong basis for causal inference, especially in a context where observational correlations would be highly suspect. The outcome measures are policy-relevant, and the comparison between mission treatment and status quo behavior is clearly motivated.

The analysis appears technically sound, with appropriate use of treatment–control comparisons and transparent reporting of effect sizes. The paper does a good job of interpreting magnitudes rather than focusing solely on statistical significance, and it appropriately benchmarks the treatment effects against baseline performance gaps. The discussion of remaining underperformance even after treatment demonstrates a nuanced understanding of what the intervention can and cannot achieve.

One potential limitation is that the paper could further explore heterogeneity in treatment effects or mechanisms (e.g., differential responses by baseline motivation or experience), though this may be constrained by data availability. Importantly, these are extensions rather than flaws in the core analysis.

Assessment: Methodologically strong and appropriate; analysis is correctly implemented and well interpreted.

Overall evaluation and score

This is a strong paper that addresses an important question using a credible and policy-relevant experimental design. It makes a meaningful contribution to the literature on public sector performance and intrinsic motivation, and it is written clearly and thoughtfully. While there is room for refinement in presentation and deeper exploration of mechanisms, the core contribution is solid and convincing.

Overall score: 85 / 100

This score reflects a paper that is clearly publishable in a good economics journal and potentially suitable for a top field journal after modest revisions.

Gemini Output

Title: Mission Motivation and Public Sector Performance Author: Yasir Khan

I. Interest, Importance, and Originality of the Research Question

The research question—whether emphasizing an organization's mission can improve the performance of *existing* public sector workers—is of high interest and policy relevance, particularly in development economics and organizational economics.

- **Importance:** Developing countries often struggle with state capacity and service delivery, partly due to the underperformance of frontline workers. While the theoretical literature has long posited that mission-oriented organizations can attract motivated workers (selection), there is a significant gap in empirical evidence regarding whether the mission can be used to motivate *incumbent* workers who are already employed under rigid civil service contracts.
- **Originality:** The paper makes a distinct contribution by moving beyond the "selection" mechanism (hiring the right people) to the "treatment" mechanism (motivating current employees). The author asserts this is the first field experiment to demonstrate this effect empirically. Furthermore, the study uniquely tests the interaction between mission motivation and financial incentives, contributing to the debate on "crowding out" vs. "crowding in" of intrinsic motivation. The setting—Lady Health Workers in Pakistan—is ideal because these workers operate in non-overlapping areas, minimizing spillovers and allowing for clean identification.

II. Clarity and Quality of Writing and Presentation

The paper is written with exceptional clarity. The experimental design, timeline, and estimation strategy are presented logically and transparently.

- **Experimental Design:** The distinction between the five treatment arms (Mission, Mission-Plus, Financial Incentive, Placebo, and Pure Control) is articulated clearly. The use of a "Placebo" group (skills refresher without mission emphasis) is a strong design choice that clarifies the results by ruling out "Hawthorne effects" or the simple benefits of extra training/attention.
- **Contextualization:** The author does an excellent job of describing the institutional setting, explaining why the lack of monitoring and the multitasking nature of the job make it a suitable testing ground for mission-based interventions.

- Visual Aids: While the text describes the randomization process well, the complexity of the five treatment arms and their sub-treatments (public/private delivery) would benefit from a visual representation.

III. Correctness and Appropriateness of the Analysis Performed

The analysis is rigorous, appropriate for a Randomized Controlled Trial (RCT), and robust to various checks.

- Methodology: The use of a stratified randomized design ensures internal validity. The primary estimation strategy (OLS with block fixed effects) is standard and appropriate.
- Data Quality: A major strength is the reliance on independent household surveys rather than worker self-reports to measure the primary outcome (home visits). This avoids social desirability bias in the outcome variable.
- Multiple Hypothesis Testing: The author correctly applies False Discovery Rate (FDR) sharpened q-values to account for multiple hypothesis testing, adding statistical rigor to the findings.
- Multitasking and Quality: The analysis goes beyond simple quantity (visits) to assess quality (multitasking). The paper convincingly shows that mission-treated workers improved on both core tasks (antenatal checks) and non-core tasks (TB screening), whereas financial incentives only increased the specific incentivized behavior. The use of Lee bounds to correct for selection bias in conditional analysis (tasks performed *conditional* on a visit) is a sophisticated and necessary robustness check.
- Mechanisms: The exploration of mechanisms is thorough. The author effectively rules out "information effects" (via the placebo) and "monitoring effects" (via survey data). The finding that mission-treated workers were more willing to work for free in a lab-in-the-field experiment provides strong evidence for the activation of altruistic preferences.
- Limitations: The author fairly acknowledges potential limitations regarding external validity, noting the specific gender of the workforce (all women) and the relatively high Human Development Index of the district. The discussion on cost-effectiveness is also valuable, showing the intervention compares favorably to other immunization programs.

IV. Overall Assessment

This is a high-quality paper that addresses a fundamental question in public economics with a clean, well-executed field experiment. It fills a critical gap in the literature by showing that management practices emphasizing mission can act as a substitute for, or complement to, financial incentives in settings with weak contract enforcement. The results are robust, the mechanisms are well-explored, and the policy implications for state capacity in developing countries are significant.

Overall Score: 94/100